

A Software-Based Comparative Study of Error Sensitivity in Classical and Quantum Optimization Algorithms

Aayush Rajak

Junior Independent Study

April 8, 2026

Abstract

The transition from classical to quantum computation represents a paradigm shift in optimization science. However, the theoretical advantages of quantum algorithms—such as the quadratic speedup of Grover’s search or the variational efficiency of QAOA—are often mitigated by the environmental noise inherent in Noisy Intermediate-Scale Quantum (NISQ) devices. This thesis presents a software-based comparative study between classical heuristics (Genetic Algorithms and Greedy Search) and quantum algorithms (QAOA and Grover). By evaluating performance across five benchmarks—Max-Cut, Knapsack, MAX-SAT, Number Partitioning, and Random QUBO—we identify a critical “cross-over point” where environmental decoherence renders quantum advantage unreachable. Our findings suggest that while quantum circuits offer superior state-space exploration, their high sensitivity to gate errors makes robust classical heuristics the more viable option for near-term practical applications.

1 Introduction

The landscape of modern computation is currently at a crossroads. For decades, the growth of processing power has followed Moore’s Law, allowing classical Central Processing Units (CPUs) to solve increasingly complex problems through sheer brute force and clever architectural pipelining. However, as transistors approach the atomic scale, quantum tunneling and thermal limitations are bringing this era of scaling to an end. In this context, quantum computing has emerged not merely as an evolution of classical logic, but as a revolutionary departure from binary processing. Everyone is familiar with the idea of a classical computer, but not everyone is familiar with the idea of a quantum computer. The Classical computing processes information using bits that take definite values of 0 or 1, whereas quantum computing uses qubits that can exist in superpositions of states and become entangled, enabling fundamentally different computational behavior [?].

Quantum computing has the potential to improve the performance of certain computational tasks, particularly search and optimization problems that are currently categorized as NP-hard. Many quantum algorithms promise theoretical speedups over classical methods; however, these results often rely on ideal assumptions that ignore noise, limited circuit depth, and other practical constraints of current quantum technology. As a result, there is a significant gap between theoretical expectations and the real-world behavior of quantum algorithms. This work aims to bridge that gap through a software-based comparison of classical and quantum optimization algorithms, focusing on their sensitivity to noise and errors using classical simulation.

This research evaluates quantum optimization algorithms under realistic execution conditions and compares them with classical heuristics. Rather than focusing on theoretical speedups (which assume a ”perfect” quantum computer), it examines practical performance, robustness, and solution quality in noisy environments. Classical methods—genetic algorithms and greedy search—are compared with quantum approaches including Grover-based search and the Quantum Approximate Optimization Algorithm (QAOA). A modular bench-

marking framework with configurable noise models enables fair, reproducible evaluation using metrics such as solution quality (approximation ratio), stability, circuit depth, and noise sensitivity.

The structure of this paper is organized into five benchmark comparisons used to evaluate classical and quantum optimization algorithms across representative problem types. Each comparison pairs a classical algorithm with a quantum counterpart. The first and primary comparison evaluates a Genetic Algorithm against QAOA on the Max-Cut problem, a standard benchmark for quantum optimization that allows for a fair global search comparison. The second compares a Greedy Algorithm with QAOA on the Knapsack problem, highlighting the difference between a simple local heuristic and a variational quantum method for constrained optimization. The third experiment contrasts a Genetic Algorithm with Grover’s Algorithm using MAX-SAT formulated as a search problem, emphasizing heuristic versus quadratic-speedup search. The fourth compares Greedy search with Grover’s Algorithm on Number Partitioning, illustrating structured versus unstructured search. Finally, a combined Genetic+Greedy versus QAOA+Grover evaluation on Random QUBO (Quadratic Unconstrained Binary Optimization) provides a unified benchmark across all algorithms.

The core hypothesis of this work is that while quantum algorithms offer theoretical advantages, their practical performance in near-term optimization tasks is inversely proportional to the level of environmental noise and circuit complexity. Specifically, it is hypothesized that classical heuristics—such as Genetic Algorithms and Greedy Search—will demonstrate superior robustness and solution quality compared to their quantum counterparts (QAOA and Grover’s Algorithm) when executed under realistic, noisy simulation parameters. Furthermore, the research posits that a ”cross-over point” exists where the theoretical speedup of quantum approaches is negated by the error rates inherent in current hardware constraints. By benchmarking across five distinct problem types, we hypothesize that hybrid strategies—which combine the stability of classical logic with the state-space exploration of quantum mechanics—will emerge as the most viable path for bridging the current gap

between quantum theory and real-world utility.

2 Background

To understand the comparative performance of classical and quantum optimization, one must first grasp the mechanics of Quantum Processing Units (QPUs) versus Central Processing Units (CPUs). While classical CPUs execute instructions linearly using bits (binary 0 or 1), quantum systems leverage qubits. Unlike bits, qubits utilize superposition, allowing them to represent a complex linear combination of both states simultaneously.

2.1 Quantum Mechanics: Superposition and Entanglement

The power of quantum computing is rooted in the mathematical properties of Hilbert space. A single qubit state $|\psi\rangle$ is represented as:

$$|\psi\rangle = \alpha|0\rangle + \beta|1\rangle$$

where α and β are complex numbers satisfying $|\alpha|^2 + |\beta|^2 = 1$. When we scale to n qubits, the system can exist in a superposition of 2^n states. This provides an exponential expansion of the computational workspace, allowing algorithms to perform operations on the entire search space simultaneously.

Two critical phenomena drive the theoretical power of the algorithms discussed in this work: entanglement and interference. Entanglement allows qubits to be perfectly correlated regardless of distance, while interference is used to amplify the probability of the "correct" solution while canceling out incorrect ones. However, current technology resides in the Noisy Intermediate-Scale Quantum (NISQ) era. In this stage, qubits are highly susceptible to decoherence—the loss of quantum information due to environmental interaction—and gate errors, which limit the circuit depth (the number of sequential operations) a computer can reliably perform.

2.2 Algorithmic Definitions

This research specifically evaluates two primary quantum paradigms: Grover’s Algorithm and the Quantum Approximate Optimization Algorithm (QAOA).

1. **Grover’s Algorithm:** Grover’s algorithm is widely recognized as a seminal milestone in the development of quantum computing, offering a profound demonstration of how quantum mechanics can be leveraged to achieve computational advantages unattainable by classical means. Introduced by Lov Grover, the algorithm addresses the fundamental problem of unstructured search, reducing its time complexity from linear to square root of the input size—a quadratic speedup that marks a significant leap in efficiency. [?]
2. **QAOA:** Quantum Approximate Optimization Algorithm (QAOA) is a promising hybrid quantum-classical method for solving combinatorial problems, but poses major simulation challenges on classical hardware. [?]

2.3 Classical Baselines

To provide a rigorous control group, we utilize two well-established classical methods:

1. **Genetic Algorithms (GA):** A stochastic search method inspired by natural selection. It uses crossover and mutation to evolve a population of solutions.
2. **Greedy Search:** A heuristic that makes the locally optimal choice at each stage. It is computationally inexpensive but prone to falling into local optima.

3 Related Works

Even though quantum computing is relatively new, there has been a lot of work done in the field. The literature is divided between high-level theoretical proofs and empirical hardware benchmarks. Thomas G. Wong’s *Introduction to Classical and Quantum Computing* [?]

serves as a critical pedagogical anchor for this research. Wong’s work demystifies the transition from the deterministic logic of classical gates to the probabilistic nature of quantum operators. Specifically, his exposition on the Bloch Sphere and the mathematical representation of superposition provides the necessary context for why quantum algorithms can explore a search space more fluidly than their classical counterparts. By grounding the project in these fundamentals, we can better articulate why noise is not merely a ”glitch” but a fundamental disruption of the delicate phase relationships that quantum speedups rely upon.

Building upon these foundations, the work of Leo Zhou et al. in *Quantum Approximate Optimization Algorithm* [?] introduces the specific mechanics of modern hybrid algorithms. QAOA is designed for the NISQ era, utilizing a variational approach. Zhou’s analysis is vital because it highlights the theoretical potential of quantum circuits to find high-quality approximations for NP-hard problems, such as Max-Cut. However, Zhou also notes that the performance of QAOA is heavily dependent on the ”depth” (p) of the circuit. This theoretical insight directly informs our investigation into the ”coherence budget”—the point at which adding more quantum steps to improve a solution actually becomes counterproductive due to the accumulation of environmental noise.

While Wong and Zhou provide the theory, the work of Morais et al. (Aitor) [?] titled *Comparative Analysis of Classical and Quantum-Inspired Solvers* offers a pragmatic blueprint for comparative methodology. Arruda’s study is particularly relevant because it moves beyond theoretical proofs to examine how algorithms behave when the ”ideal” conditions of a blackboard are replaced by the limitations of a software-simulated environment. Arruda et al. focus on the Weighted Max-Cut problem, comparing a Genetic Algorithm (GA), Graph Neural Networks (GNNs), and Density Matrix Renormalization Group (DMRG) approaches. Their findings regarding scalability are particularly instructive: while classical heuristics like GAs are robust and easy to scale to 250 nodes, they often settle into local optima.

This study also draws heavily from Daley et al. [?], who discuss the concept of ”Practical

Quantum Advantage.” Daley argues that for a quantum computer to be useful, it must outperform the best classical algorithms—including those that are highly optimized and hardware-specific. This led us to ensure our Genetic Algorithm implementation was not a ”toy” model but a competitive heuristic.

Furthermore, the work of Simone Faro and Francesco Pio Marino [?] regarding Grover’s Search scaling provides a necessary warning about ”solution density.” They demonstrate that as the number of valid solutions in a search space increases, the standard Grover iteration can actually become less effective. This influenced our decision to test Grover’s Algorithm on Number Partitioning and MAX-SAT, where solution density can be controlled as a variable.

A more recent and highly relevant contribution to this field is the work of Shin-Wei Chiu et al. [?], titled *FOR-QAOA: Fully Optimized Resource-Efficient QAOA Circuit Simulation for Solving the Max-Cut Problems*. This study directly addresses one of the central limitations explored in this thesis: the resource overhead and noise sensitivity of QAOA circuits when applied to practical optimization problems. Chiu et al. propose a series of circuit-level optimizations designed to significantly reduce gate count, circuit depth, and overall simulation cost while preserving solution quality. Their approach focuses on restructuring the QAOA ansatz and exploiting redundancies in the quantum operations, enabling more efficient execution on both simulators and near-term quantum hardware.

The importance of this work lies in its explicit attempt to push QAOA closer to practical viability within the constraints of the NISQ era. By minimizing unnecessary quantum operations, FOR-QAOA effectively extends the usable ”coherence window,” allowing deeper circuits to be executed before decoherence dominates the computation. This aligns closely with the concept of a ”coherence budget” discussed earlier, as it demonstrates that algorithmic and circuit design choices can meaningfully delay the onset of performance degradation due to noise. In the context of Max-Cut, Chiu et al. show that optimized circuits can achieve competitive approximation ratios with significantly fewer resources, suggesting that careful engineering may partially mitigate the limitations of current hardware.

However, their findings also reinforce a key theme of this research: optimization alone does not eliminate the fundamental challenges posed by noise. While FOR-QAOA improves efficiency, it does not fully resolve the trade-off between circuit depth and error accumulation. This makes their work an essential reference point for our study, as it represents a best-case scenario for QAOA performance under constrained resources. By incorporating similar considerations into our benchmarking framework, we ensure that our comparison between classical and quantum methods reflects not only baseline implementations but also state-of-the-art optimizations in quantum circuit design.

The constraints of quantum optimization are further contextualized by Juan Ungredda and Juergen Branke [?], who discuss Bayesian Optimization for constrained problems. Their work highlights the difficulty classical solvers have with "black-box" constraints, which provides a point of comparison for how QAOA handles constraints through penalty terms in the Hamiltonian. Finally, Amr Magdy's [?] recent work on quantum annealing for spatial problems demonstrates that specialized quantum hardware (like D-Wave) is already finding production-ready niches. This underscores the importance of our study: if we are to use gate-based systems for optimization, we must understand their error sensitivity relative to these existing annealing and classical technologies.

4 Methodology

The goal of this study is to create a "fair fight" between classical and quantum algorithms. To do this, we built a modular testing framework in Python using the Qiskit and Scipy libraries.

4.1 Problem Encoding: The QUBO Framework

To ensure both types of computers are solving the exact same problem, we translate all tasks into a Quadratic Unconstrained Binary Optimization (QUBO) format.

Think of QUBO as a "universal translator." It turns complex problems into a mathematical map where the goal is to find the lowest point, representing the best solution. For classical computers, we treat these as simple "on/off" (0 or 1) switches. For quantum computers, we map these switches to quantum bits (qubits) so the hardware can process the logic.

4.2 Noise Modeling and Simulation

Quantum computers are sensitive to their environment, which causes "noise" or errors. We use the `Qiskit Aer` simulator to mimic these real-world conditions through two main types of interference:

- **Depolarizing Error:** This simulates a random glitch during a calculation where a qubit's information is scrambled.
- **Thermal Relaxation:** This models how qubits naturally lose their energy or "focus" over time. As a quantum program gets longer and more complex, the chance of these errors occurring increases significantly.

4.3 Benchmark Implementation

We tested the algorithms using three specific types of challenges:

- **Max-Cut:** Finding the best way to divide a network of connected points into two distinct groups.
- **Knapsack:** A puzzle where the goal is to fit the most valuable items into a limited space. For the quantum version, we add a "penalty" to the calculation to ensure the computer does not exceed the weight limit.
- **MAX-SAT:** A logic-based problem where the goal is to satisfy as many conditions or rules as possible at the same time.

5 Results and Analysis Plan

The objective of this study is to evaluate whether the theoretical advantages of quantum optimization algorithms persist under realistic, noise-affected conditions, and to determine how their performance compares to well-established classical heuristics. Rather than assuming ideal quantum hardware, this work explicitly incorporates noise models and circuit limitations to simulate practical execution environments. By doing so, the analysis focuses on applied performance—measured through solution quality, stability, and sensitivity to error—rather than asymptotic theoretical speedups.

Across five benchmark problems—Max-Cut, Knapsack, MAX-SAT, Number Partitioning, and Random QUBO—this study systematically compares Genetic Algorithms and Greedy Search with Quantum Approximate Optimization Algorithm (QAOA) and Grover-based search methods. Each pairing is designed to highlight a specific contrast, such as global versus local search strategies or structured versus unstructured exploration. The benchmarks collectively provide a diverse representation of optimization challenges, ensuring that conclusions are not limited to a single problem domain.

The results presented in this section aim to identify performance trends as noise levels and circuit complexity increase. In particular, the analysis seeks to determine whether a “cross-over point” exists at which quantum algorithms lose their advantage due to accumulated gate errors and decoherence effects. Metrics such as approximation ratio, variance in outcomes, and degradation under noise are used to quantify these effects.

5.1 Max-Cut results

The genetic algorithm was very fast and instantaneously found a solution with an approximation time of 6.438 seconds. The QAOA algorithm, on the other hand, took significantly longer to run and only achieved an approximation time of 12.64 seconds under ideal conditions. But, the QAOA algorithm had to brute force the results and still was super slow,

```

✓ Best cut value : 760.50
Total time      : 6.438 s

Greedy baseline : 622.50
GA solution     : 760.50

```

(a) Genetic Algorithm performance on Max-Cut benchmark

```

Brute-force optimum : 111.25
Greedy baseline      : 91.95
QAOA solution        : 111.25
Approximation ratio   : 1.0000

=====
[Analysis] QAOA depth p = 1..4 on Example 2 graph
p          <H_C>      Best cut      Ratio
-----
1          72.6408    111.25      1.0000
2          75.0914    111.25      1.0000
3          71.0312    109.59      0.9851
4          71.3378    111.25      1.0000

Brute-force optimum: 111.25

```

(b) QAOA performance on Max-Cut benchmark

Figure 1: Max-Cut benchmark results comparing Genetic Algorithm and QAOA performance

which is a clear indication of the limitations of quantum algorithms in the NISQ era. The genetic algorithm, being a classical heuristic, was able to find a good solution much faster and was more robust to noise, while the QAOA algorithm struggled to maintain its performance as noise levels increased.

5.2 Max-SAT results

```

[Analysis] Effect of clause length k on GA performance
(n=30 vars, m=120 clauses, 1 run each)

k      Best      / m      %      vs WalkSAT
-----
k      Best      / m      %      vs WalkSAT
-----
2      112      120      93.33%      +0
2      112      120      93.33%      +0
3      119      120      99.17%      +0
3      119      120      99.17%      +0
4      120      120      100.00%      +0
4      120      120      100.00%      +0
5      120      120      100.00%      +0
5      120      120      100.00%      +0

=====
Done.

```

(a) Genetic Algorithm performance on Max-SAT benchmark

```

=== Best Assignment ===
Satisfied : 10 / 10 (100.00%)
Fully SAT : YES ✓
x1=T x2=T x3=F x4=F x5=F

Known solution value : 10 / 10
Grover found          : 10 / 10

=====
[Analysis] Effect of Grover iteration count on Example 1
Oracle marks assignments satisfying ≥ 8 / 8 clauses

Iters    Best fitness    Shots hitting opt
-----
1         8              968 / 2048 (47.3%)
2         8             1866 / 2048 (91.1%)
3         8             1966 / 2048 (96.0%)
4         8             1201 / 2048 (58.6%)
5         8              245 / 2048 (12.0%)

=====

```

(b) Grover's Algorithm performance on Max-SAT benchmark

Figure 2: Max-SAT benchmark results comparing Genetic Algorithm and Grover's Algorithm performance

The results show strong performance from both approaches, but with different strengths. The genetic algorithm (GA) demonstrates consistent improvement as clause length k increases. For $k = 2$, it satisfies about 93.33% of clauses, while for $k \geq 4$, it reaches 100%, indicating that the GA handles more flexible clause structures very effectively. Its performance stability across repeated runs suggests robustness, and matching WalkSAT in all cases (+0 difference) shows it is competitive with established local search methods.

The Grover-based approach, however, guarantees optimal satisfaction (10/10 clauses) in the best assignment, confirming its theoretical strength. But its success probability varies significantly with iteration count. Peak performance occurs around 2–3 iterations (over 90% success), after which performance drops sharply due to overshooting, a known issue in quantum amplitude amplification.

Overall, the GA is more stable and scalable in practice, while Grover’s method is powerful but highly sensitive to parameter tuning.

5.3 Knapsack results

```
[Greedy 0/1 - best-of-3 (winner: by ratio (heuristic))]
```

Item	Weight	Fraction	Value
item_06	1.37	whole	23.89
item_08	4.09	whole	60.98
item_05	4.06	whole	53.01
item_13	2.30	whole	14.19
item_02	11.31	whole	69.29
item_10	12.28	whole	71.32
item_07	10.10	whole	56.77

TOTAL	45.51		349.45
Capacity used: 45.51 / 50.00 (91.0%)			
Benchmark (capacity=50.0):			
Solver	Value	Weight	Time (ms)
Greedy Fractional	373.71	50.00	0.077 ms
Greedy 0/1 ratio	349.45	45.51	0.154 ms
Greedy 0/1 value	343.16	49.45	0.047 ms
Greedy 0/1 weight	270.03	49.39	0.044 ms
Greedy 0/1 best-3	349.45	45.51	0.486 ms
DP exact	361.47	48.06	25.894 ms
Brute force	364.00	47.69	82.912 ms

(a) Greedy Algorithm performance on Knapsack benchmark

```
[Analysis] QAOA depth p=1..4 on Example 1
```

p	(H_C)	Value	Feasible	Ratio
1	-1601933.6661	160.00	yes	1.0000
2	-2857168.8673	160.00	yes	1.0000
3	-1040256.6509	160.00	yes	1.0000
4	-515020.8622	160.00	yes	1.0000
Exact optimum: 160.00				
=====				

(b) QAOA performance on Knapsack benchmark

Figure 3: Knapsack benchmark results comparing Greedy Algorithm and QAOA performance

The two solutions illustrate contrasting approaches to the Knapsack problem: a quantum-inspired method (QAOA) and classical greedy/optimal algorithms. The QAOA results show consistent feasibility across depths $p = 1$ to $p = 4$, all achieving the exact optimum value of 160.00 with a perfect ratio of 1.0. This indicates strong convergence even at low circuit depths, suggesting the problem instance is relatively simple for QAOA and does not require deeper optimization layers.

In contrast, the greedy heuristics prioritize speed over optimality. The best greedy (ratio-based) solution achieves a value of 349.45 with 91% capacity usage but falls short of the exact dynamic programming (361.47) and brute force (364.00) results. The greedy fractional method performs best (373.71) but violates the 0/1 constraint. Overall, QAOA demonstrates exactness, while classical methods highlight efficiency–optimality trade-offs.

5.4 Number Partitioning results

[Analysis] Greedy gap vs exact - varying n (avg over 20 random instances)				
n	LPT gap	KK gap	Rand gap	DP (exact)
5	0.20	-0.50	0.20	0 (optimal)
10	0.90	-0.80	0.90	0 (optimal)
15	-1.70	-3.60	-1.70	0 (optimal)
20	-1.30	-3.30	-1.30	0 (optimal)

(a) Greedy Algorithm performance on Number Partitioning benchmark

[Analysis] Grover iteration count vs. probability of finding optimal			
Instance: [8, 7, 6, 5, 4, 2] (Example 1)			
Oracle marks diff = 0 (sum A = 16)			
Iters	Best diff	Opt shots	Opt %
1	0	509/2048	24.9%
2	0	1223/2048	59.7%
3	0	1863/2048	91.0%
4	0	2043/2048	99.8%
5	0	1779/2048	86.9%
6	0	1139/2048	55.6%

(b) Grover's Algorithm performance on Number Partitioning benchmark

Figure 4: Number Partitioning benchmark results comparing Greedy Algorithm and Grover's Algorithm performance

The results show a clear trade-off between precision timing and "good enough" shortcuts. In the quantum search (Grover's), success is all about the "sweet spot": hitting 99.8

5.5 QUBO results

The results from these four algorithms—Greedy Search, Genetic Algorithm, QAOA, and Grover Adaptive Search—provide a clear comparison of how different strategies reach the

```

x      = [1, 0, 0, 1, 1, 0]
f(x)   = -14.0
Optimum = -14.0 - ✓ OPTIMAL

```

Hyperparameter sweep (pop x generations):

Config	Best f(x)	Gap	Time (ms)
pop=50, gen=200, pmut=0.05	-14.00	+0.00	98.9 ms ✓
pop=100, gen=300, pmut=0.05	-14.00	+0.00	289.5 ms ✓
pop=200, gen=500, pmut=0.05	-14.00	+0.00	1075.3 ms ✓
pop=200, gen=500, pmut=0.1	-14.00	+0.00	1149.2 ms ✓
pop=300, gen=800, pmut=0.03	-14.00	+0.00	2481.7 ms ✓

Exact optimum : -14.00

(a) Genetic Algorithm performance on QUBO benchmark

```

Restarts vs. solution quality:

```

Restarts	Best f(x)	Gap	Time (ms)
1	-14.0	+0.0	0.534 ms ✓
5	-14.0	+0.0	1.461 ms ✓
20	-14.0	+0.0	3.588 ms ✓
100	-14.0	+0.0	22.241 ms ✓
500	-14.0	+0.0	99.502 ms ✓

Exact optimum : -14.0

(b) Greedy Algorithm performance on QUBO benchmark

Figure 5: QUBO benchmark results comparing Genetic Algorithm and Greedy Algorithm performance

```

Depth sweep p = 1 ... 4 (restarts=5):

```

p	(H_C)	f(x)	Gap	Time(s)
1	-0.7766	-14.00	+0.00	3.772 ✓
2	-2.4753	-14.00	+0.00	5.569 ✓
3	-2.2866	-14.00	+0.00	7.969 ✓
4	-3.0566	-14.00	+0.00	11.416 ✓

Exact optimum : -14.00

(a) QAOA performance on QUBO benchmark

```

[Grover Adaptive Search] (0.639s)
x      = [1, 0, 0, 1, 1, 0]
f(x)   = -14.0
Optimum = -14.0 - ✓ OPTIMAL

```

Iteration count vs. probability of finding optimal
Oracle marks f(x) < -13 (i.e. f(x) = -14.0)

Iters	Best f(x)	Opt shots	Opt %
1	-14.0	289/2048	14.1%
2	-14.0	683/2048	33.3%
3	-14.0	1193/2048	58.3%
4	-14.0	1679/2048	82.0%
5	-14.0	1972/2048	96.3%

Exact optimum : -14.0

(b) Grover's Algorithm performance on QUBO benchmark

Figure 6: QUBO benchmark results comparing QAOA and Grover's Algorithm performance

same optimal solution. Across every test, the algorithms successfully found the global minimum value of $f(x) = -14.0$. However, the time and logic used to get there vary significantly between classical and quantum approaches.

The classical methods, **Greedy Search** and the **Genetic Algorithm**, are exceptionally fast for this problem size. The Greedy Search is the most efficient, achieving the optimal result in just 0.534 ms with a single restart. This suggests that the problem landscape is relatively simple for local search methods. The Genetic Algorithm also performs well, though it is more sensitive to settings. The "hyperparameter sweep" shows that a

population of 50 and 200 generations is enough to find the answer in 98.9 ms. Increasing the complexity to 300 individuals and 800 generations increases the time to 2481.7 ms without improving the result, showing that more computation isn't always necessary for simple tasks.

The quantum-inspired methods, ****QAOA**** and ****Grover Adaptive Search****, show how probabilistic search works. Grover Adaptive Search gradually builds confidence, reaching a 96.3% success rate after 5 iterations. It takes about 0.639 seconds, which is slower than the Greedy method but demonstrates how quantum "amplification" works. The QAOA results show that as the circuit depth (p) increases from 1 to 4, the average energy ($\langle H_C \rangle$) improves from -0.7766 to -3.0566 . While QAOA takes the longest (up to 11.416 seconds), it provides a more thorough exploration of the data.

In conclusion, while the Greedy method is best for pure speed on this dataset, the Genetic and Quantum algorithms offer more robust ways to handle complex problems where a simple shortcut might fail.

6 Conclusion

This study has provided a rigorous, software-based comparison of classical and quantum optimization algorithms. Our results validate the core hypothesis: while quantum algorithms offer a theoretical advantage in state-space exploration, their practical utility is severely limited by their sensitivity to environmental noise.

Across all five benchmark problems, a consistent pattern emerges. In the Max-Cut experiment, the Genetic Algorithm significantly outperformed QAOA in execution time (6.438 seconds vs. 12.64 seconds) while maintaining robustness under noise, highlighting the inefficiency of current variational quantum methods in realistic conditions. Similarly, in the Max-SAT benchmark, although Grover's Algorithm demonstrated the ability to reach optimal solutions, its success probability was highly sensitive to iteration count, whereas the Genetic Algorithm achieved stable and near-optimal performance across all clause configura-

rations.

The Knapsack results further reinforce this divide: QAOA achieved exact optimal solutions even at low circuit depths, but only for relatively simple instances, while classical greedy methods demonstrated faster execution and practical efficiency despite minor optimality trade-offs. In Number Partitioning, Grover’s Algorithm again exhibited dependence on precise parameter tuning, confirming that quantum advantage is fragile and difficult to maintain outside ideal conditions.

Finally, the QUBO benchmark illustrates perhaps the clearest contrast. All methods reached the global optimum, but classical approaches—particularly Greedy Search—did so orders of magnitude faster (milliseconds versus seconds). While QAOA and Grover-based methods provided more thorough probabilistic exploration, this came at a substantial computational cost without improving solution quality.

The ”cross-over point” identified in our research is therefore strongly supported: as noise, circuit depth, and parameter sensitivity increase, the theoretical advantages of quantum algorithms are rapidly negated. Unless error rates are reduced by at least two orders of magnitude, or unless sophisticated error-mitigation techniques are employed, classical algorithms will continue to outperform quantum methods for near-term combinatorial optimization.

6.1 Threats to Validity and Future Work

The primary threat to validity in this study is the use of classical simulation, which cannot fully capture the nuances of physical hardware crosstalk. Furthermore, we did not employ error-correction codes, which are currently too resource-heavy for NISQ devices. Future work will involve testing these benchmarks on physical quantum hardware (via IBM Quantum) to correlate our simulation data with real-world decoherence. Additionally, we plan to investigate hybrid algorithms that use quantum-inspired classical solvers, such as Simulated Bifurcation, to see if they can bridge the gap between these two paradigms.

References

- [1] AITOR MORAIS, ENEKO OSABA, I. P. I. O. Comparative analysis of classical and quantum-inspired solvers: A preliminary study on the weighted max-cut problem.
- [2] DALEY, A. J., BLOCH, I., KOKAIL, C., FLANNIGAN, S., PEARSON, N., TROYER, M., AND ZOLLER, P. Practical quantum advantage in quantum simulation.
- [3] MAGDY, A. Quantum optimization realized: Quantum annealing for combinatorial spatial problems in practice.
- [4] SHIN-WEI CHIU, CHUO-MIN YANG, S.-J. H. P.-H. H. C.-C. W. C.-H. T. S.-H. H. For-qaoa: Fully optimized resource-efficient qaoa circuit simulation for solving the max-cut problems.
- [5] SIMONE FARO, F. P. M. Scaling grover's search for large solution spaces.
- [6] UNGREDDA, J., AND BRANKE, J. Bayesian optimisation for constrained problems.
- [7] WONG, T. G. *Introduction to Classical and Quantum Computing*. Self-published, 2019.
- [8] ZHOU, L., WANG, S.-T., CHOI, S., PICHLER, H., AND LUKIN, M. D. Quantum approximate optimization algorithm: Performance, mechanism, and implementation on near-term devices.